

BSc (Hons) in Information Technology Specializing Data Science

Research Project - IT4010

MLOPs Report

Group ID: 25-26J-103

Project Title: E-Learning Platform for hearing impaired Students

1. IT22546852 : AI-DRIVEN REAL-TIME SLSL E-LEARNING WITH GESTURE CORRECTION & 3D AVATAR STYLES
2. IT22143754: AI-POWERED SINHALA LIP-READING ASSISTANT WITH 3D AVATAR FEEDBACK
3. IT22255488: EMOTION AWARE MUSIC FOR HEARING IMPAIRED STUDENTS
4. IT22561466: EMOTION MONITORING WITH REAL-TIME FEEDBACK AND STAGEBASED LEARNING OPTIMIZATION

		IT22546852	IT22143754	IT22255488	IT22561466
Data Pipeline	Data Sources:	- SLSL dataset from kaggle - Custom Sri Lankan Sign Language (SSL) gesture dataset (images/videos collected from hearing-impaired individuals and online SSL resources) - Annotated gesture samples labeled according to SSL vocabulary	custom Sinhala lip-reading dataset (Participants perform predefined words and sentences. Each word and sentence is recorded using a standard webcam setup ensuring consistent background, lighting, and camera positioning.	Audio files (WAV, MP3, FLAC, OGG, M4A) labeled with valence & arousal values on a 1–9 scale sourced from music emotion datasets	- Custom facial expression and attention monitoring dataset collected using webcam video recordings - Public face detection and emotion recognition datasets (e.g., FER2013 / MediaPipe sample facial datasets)

BSc (Hons) in Information Technology Specializing Data Science

Research Project - IT4010

MLOPs Report

					• Custom dataset containing student face images and short video clips for gaze detection and screen attention monitoring • Annotated labels for: ✓ focused ✓ distracted ✓ confused ✓ absent ✓ looking at screen ✓ looking away
	Data Preprocessing:	• Frame extraction from gesture videos • Hand landmark detection using MediaPipe • Normalization of hand coordinates (scaling, centering) • Data augmentation (rotation, flipping, noise injection) to improve robustness	• Face detection • Lip region extraction (ROI) • Landmark detection (MediaPipe Face Mesh) • Frame normalization (resize, scaling)	• Resample audio to 44,100 Hz mono • Segment into 5-second chunks (220,500 samples); zero-pad partials > 2.5 s • Compute 128-band mel-spectrogram (n_fft=2048, hop_length=512) • Convert to dB scale via librosa.power_to_db() • MinMax-normalise valence & arousal labels (scikit-learn) Trained model serialised	• Face detection using OpenCV Haar Cascade / MediaPipe Face Detection • Facial landmark extraction using MediaPipe Face Mesh • Eye region detection for gaze tracking

BSc (Hons) in Information Technology Specializing Data Science

Research Project - IT4010

MLOPs Report

		Sequence padding/truncation for consistent input length	Sequence creation for temporal frames	as best_cnn_model.keras (Keras format); loaded at server start-up for inference. Training label statistics (min/max) stored for denormalisation.	- Head pose estimation for screen attention monitoring - Face alignment and normalization - Image resizing and pixel scaling - Frame extraction from webcam video stream - Sequence generation for temporal attention analysis - Data augmentation: ✓ brightness adjustment ✓ rotation ✓ slight zoom ✓ horizontal shift - Label encoding for attention states
	Data storage and versioning	Dataset stored in structured directories	- Dataset stored as: - Frame sequences	Trained model serialised as best_cnn_model.keras (Keras format); loaded at server start-	Face image datasets stored in train /

BSc (Hons) in Information Technology Specializing Data Science

Research Project - IT4010

MLOPs Report

		(train/validation/test split) - Version control using Git - Preprocessed data stored as NumPy for faster training	- Labels (word + phoneme) - Version control is maintained using Git-based tracking	up for inference. Training label statistics (min/max) stored for denormalisation.	validation / test directory structure - Preprocessed facial frames stored as NumPy arrays (.npy) for efficient model training - Video frame sequences stored with timestamp-based naming - Version control maintained using Git - Model experiment tracking using version-based folders - Dataset versions maintained for reproducibility
Model Development	Model Selection	- CNN + LSTM hybrid architecture - CNN → spatial feature extraction (hand shape)	- CNN → extracts lip shape - LSTM → captures temporal motion	CNN Convolutional Neural Network mapping mel-spectrograms to Russell's Circumplex Model of Affect (2-D valence/arousal space). Architecture includes batch	CNN architecture for face feature extraction CNN + LSTM hybrid model for temporal

BSc (Hons) in Information Technology Specializing Data Science

Research Project - IT4010

MLOPs Report

		LSTM → temporal sequence learning (gesture motion)		normalisation and dropout layers.	attention monitoring - CNN extracts: ✓ facial features ✓ eye region features ✓ expression patterns - LSTM captures: ✓ gaze movement sequence ✓ head movement over time ✓ temporal attention changes
	Model Training	- Loss Function: Categorical Cross-Entropy - Optimizer: Adam - Evaluation Metrics: Accuracy, Precision, Recall - Training done using TensorFlow - Early stopping and dropout used to prevent overfitting - Model Integration	Training pipeline includes: - Input: Lip frame sequences - Output: Word / viseme class Config: - Loss: Categorical Cross-Entropy - Optimizer: Adam - Metrics: Accuracy Training pipeline automated Hyperparameters tracked Model versions saved	- Framework: TensorFlow / Keras - Loss: MSE Optimiser: Adam - Labels: MinMax-normalised to [0,1]; denormalised post-prediction - Valence range: [1.6, 8.4] · - Arousal range: [1.8, 8.1] - Evaluation: Multi-run variance testing for prediction consistency	- Loss Function: Categorical Cross-Entropy - Optimizer: Adam - Metrics: ✓ Accuracy ✓ Precision ✓ Recall ✓ F1-score - Early stopping used to reduce overfitting

BSc (Hons) in Information Technology Specializing Data Science

Research Project - IT4010

MLOPs Report

					- Dropout layers added for regularization - Training performed using TensorFlow / Keras - Hyperparameters tracked: ✓ learning rate ✓ batch size ✓ epochs ✓ dropout rate
Model Integration	Tools used	- Frontend: React (for interactive learning UI) - Backend: Flask REST API - ML Framework: TensorFlow - Computer Vision: MediaPipe (hand tracking) + OpenCV - Real-time Processing: WebRTC / browser webcam integration Model Deployment	- Frontend: React (interactive learning interface) - Backend: Flask REST API - Machine Learning: TensorFlow - Computer Vision: MediaPipe + OpenCV	TensorFlow ≥ 2.20 / Keras — inference librosa ≥ 0.10 — mel-spectrogram & audio resampling NumPy ≥ 1.26 — array operations scikit-learn ≥ 1.2 — MinMaxScaler normalization FastAPI ≥ 0.104 + uvicorn — REST endpoint /predict soundfile ≥ 0.12 — audio I/O python-multipart ≥ 0.0.6 — WAV blob upload handling	- Frontend: React - Backend: Flask REST API - ML Framework: TensorFlow / Keras - Computer Vision: ✓ OpenCV ✓ MediaPipe Face Mesh - Real-time webcam integration: WebRTC

BSc (Hons) in Information Technology Specializing Data Science

Research Project - IT4010

MLOPs Report

Model Deployment	Testing Environments	- Local development (localhost with webcam testing) - Controlled user testing with sample SSL learners - Cross-browser testing (Chrome preferred for MediaPipe stability)	- Local (developer testing) - Staging (integration testing) - User testing (hearing-impaired learners) Testing includes: - Lip detection accuracy - Prediction correctness - Feedback clarity	Local development server; API health monitored via /health endpoint polled every 3 s from the React frontend. Multi-run variance testing used to validate prediction consistency.	- Controlled testing with sample users - Different lighting condition testing - Different face angle testing - Browser-based testing - Real-time latency testing Testing includes: - face detection accuracy - gaze tracking correctness - attention state classification - alert trigger validation
	Deployment Platform	- Localhost deployment for development - Flask API deployment on cloud server - Frontend deployment using Vercel	- Localhost deployment for development - Flask API deployment on cloud server - Frontend deployment using Vercel	Python FastAPI server on localhost port 8000; served via uvicorn ASGI. CORS middleware enabled for browser-to-localhost communication.	- Localhost deployment for development - Flask API deployment on cloud server - Frontend deployment using Vercel

BSc (Hons) in Information Technology Specializing Data Science

Research Project - IT4010

MLOPs Report

		Backend deployment using Railway	Backend deployment using Railway		Backend deployment using Railway
	Deployment Method	- Model serialized as .h5 / SavedModel - API endpoint for real-time gesture prediction - Continuous inference from webcam stream - Feedback loop integrated into UI	- Train model - Save version - Deploy API - Serve predictions	Model loaded from .keras file at server startup. Frontend polls /predict every 500 ms with a 2-second WAV chunk (encoded in-browser); server responds in < 200 ms for near real-time emotion tracking.	- Model serialized as .h5 / SavedModel - REST API endpoint for real-time prediction - Continuous inference from webcam stream - Automated prediction feedback to frontend - Version-based model deployment
Future Enhancements	Model Improvement:	- Expand SSL dataset for better accuracy and regional variation coverage - Improve real-time feedback with confidence scoring and correction hints - Integrate 3D avatar-based gesture	- Expand Sinhala lip dataset - Improve phoneme-to-viseme mapping ☐ Add adaptive feedback system	- Move beyond 4-quadrant categorisation to the full continuous valence-arousal space with nuanced emotion labels (e.g. nostalgia, triumph, serenity) - Offline-first architecture: package the CNN model with a local speech recognition engine	- improve face detection robustness under low light - advanced gaze estimation model - head pose-based attention scoring

BSc (Hons) in Information Technology Specializing Data Science

Research Project - IT4010

MLOPs Report

		demonstration for better learning experience		(e.g. Whisper) to remove cloud dependencies Expanded training data for improved accuracy across diverse music genres	- personalized focus behavior model - multi-face classroom support - emotion + attention fusion model
--	--	--	--	--	---